Udemy- Data Science

Machine Learning

- Data Acquisition
- Data Cleaning
- Train, Val (For adjusting Hyperparameters) & Test Data
- Model Training and Building
- Model Testing
- Adjust Model parameters
- Deploy models

Model Evaluation

1. **Key Classification Metric**
    a. **Accuracy**
        i. correct / total
        ii. Useful when target classes are well balanced
    b. **Recall**
        i. TP / (TP + FN)
        ii. Ability of a model to predict all relevant cases
    c. **Precision**
        i. TP / (TP + FP)
        ii. Ability of a model to predict only relevant cases.
    d. **F1- Score**
        i. 2 * (P * R) / (P + R)
        ii. Combination of Precision and Recall, Harmonic Mean.
        iii. Harmonic Mean, HM punishes extreme values
    e. **Confusion Matrix**
        i. **TP FN**
        ii. **FP TN**

2. **Key Regression Metric**
    a. **Mean Absolute Error**
    b. **Mean Squared Error**
        i. Punishes extreme values
    c. **Root Mean Squared (RMS)**
        i. Square root of MSE
        ii. Most popular, has same unit of ‘y’

Scikit Learn

1. From sklearn.family import Model (Estimator)
2. Split Train & Test data
3. Model.fit(X_train, y_train)
4. Model.predict(X_train)
5. Model.score()

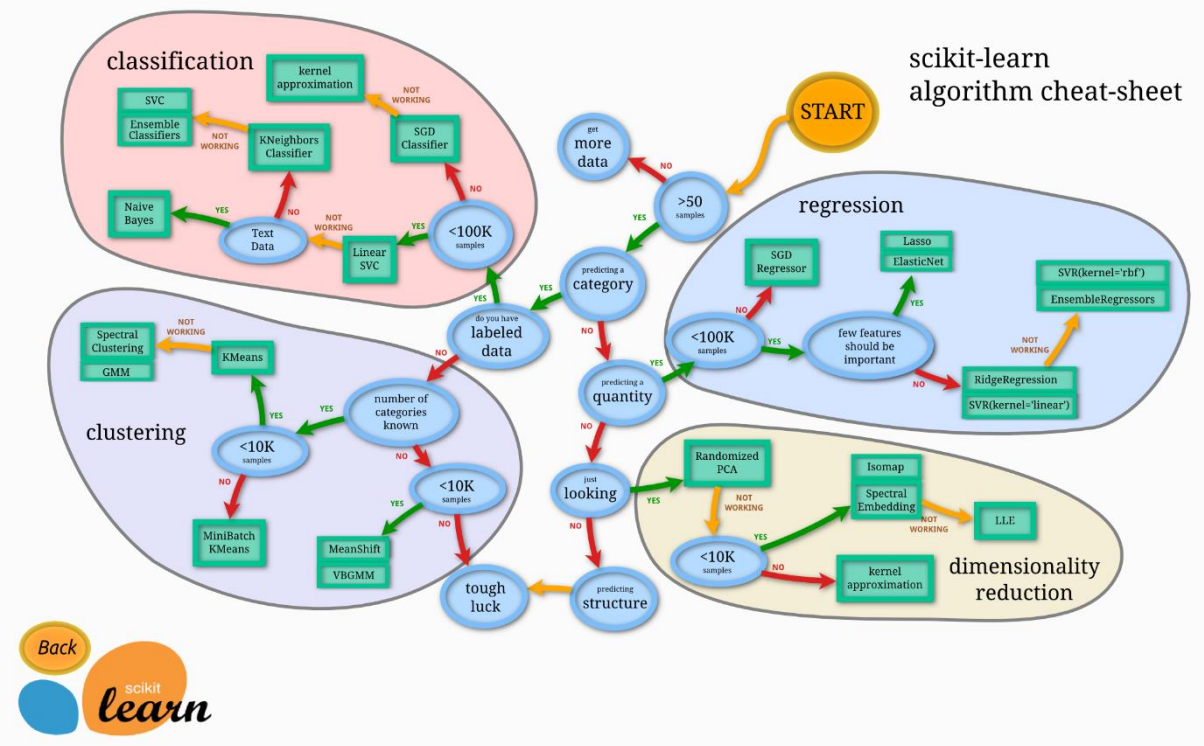

Bias Variance Trade off

- Point where we are just adding noise by adding model complexity
- Training error goes down, but test error is starting to go up
- After this point, model begins to overfit

Logistic Regression

- Sigmoid / Logistic Function can take values from 0 & 1
- Confusion Matrix can be used for model evaluation.

KNN

- Classification algorithm
- High prediction cost
- Not for high dimension
- Not for categorical features

Decision Tree

- **Nodes,** split for the value of a certain attribute
- **Edges,** outcome of a split to next node
- Entropy and Information Gain are the methods to choose best split

Random Forest (Ensemble)

- Many trees with a random sample of features chosen as split.
- **Advantage**
    - If a very strong feature in data set, then most of the bagged trees will use that feature as top split, resulting in highly corelated ensemble.

- By randomly leaving out candidate features from each split, Random Forest decorrelate the trees.
- Averaging then will reduce the variance.

Support Vector Machines (SVM)

- Choose a hyperplane which maximise margin between classes
- Expand this idea to non-linear space using kernel trick

K Means

- Unlabelled data in Unsupervised Learning
- Group similar clusters together
  - Compute cluster centroid by taking mean vector of points.
  - Assign each data point to the cluster for which the centroid is closest.

Principal Component Analysis (PCA)

- General factor analysis
- Unsupervised
- Find which feature explains most variance in data

Recommendation Systems

1. **Content Based**
   - **a.** Focus on attributes of items
   - **b.** Based on similarity between items
2. **Collaborative Filtering**
   - **a.** Wisdom of crowd
   - **b.** Based on the knowledge of user's attitude to items
   - **c.** More commonly used, better result

d. Memory Based or Model Based

Natural Language Processing (NLP)

- **Bag of Words,** A document represented as a vector of words.
- Cosine similarity on the vectors to determine similarity.
    - sim (A, B) = cos(theta) = A.B / (|A| * |B|)
- Improve in Bag of words by adjusting word counts based on their frequency in corpus.
- We can use **TF / IDF**
    - **TF (Term Frequency)**
        - Importance of term within document
        - Number of occurrences of term in the document
    - **IDF (Inverse Document Frequency)**
        - Importance of term in the corpus
        - IDF(t) = log(D/t)
        - log (Total Documents / no. of documents with term)
    - **TF-IDF = TF * IDF**

Big Data

Distributed System, which distribute data over multiple machines.

- Easier to scale

Hadoop is a way to distribute very large files across multiple machines.

- **HDFS, Hadoop Distributed File System**
    - Allows to work with large dataset
    - Duplicate blocks for fault tolerance
    - **MapReduce** allows computation on data
    - Has a name node & various data nodes attached.

- Uses blocks of data of 128MB default each replicated 3 times.
- Blocks are distributed in a way to support fault tolerance.
- **Map Reduce** is a way to split a computation task to a distributed set of files such as HDFS.
- It consists of Job Trackers and Task Trackers.
- Job trackers send code to run on task tracker.
- Task trackers allocate memory and CPU and monitor the task on worker node.

Spark

- Quickly and easily handle big data.
- Open Source on Apache
- Flexible alternate to MapReduce
- Spark can use data stored in a variety of formats
  - Cassandra, S3, HDFS … etc

Spark vs MapReduce

- Requires files to be stored in HDFS, Spark does not
- Spark can perform 100x faster
  - Spark keeps all the data in memory after each transformation, MapReduce only to disk.
  - Spark can use disk if memory is full

At the core RDD, Resilient Distributed Dataset

- RDD has 4 main features
  1. Distributed collection of data
  2. Fault Tolerant
  3. Parallel operation – partitioned
  4. Ability to use many data sources
- RDDs are immutable, lazily evaluated, and cacheable

- 2 types of RDD operations
    - Transformations
        - Filter, apply filter and return elements that are true
        - Map, transform each element
        - FlatMap, transform each element to 0-N elements.
    - Actions
        - Collect, return all elements of RDD as an array
        - Count, number of elements
        - First, first element in RDD
        - Take, return array with first n element

- Offer RDD is holding values in tuples (key, value)
    - Reduce, aggregate RDD elements using function that return single value
    - ReduceByKey,
        - Similar to Group by